

```
import pandas as pd
import numpy as np

dataset = pd.read_csv("Material Compressive Strength Experimental Data.csv")
```

```
dataset.head(5)
```

	Material Quantity (gm)	Additive Catalyst (gm)	Ash Component (gm)
0	486.42	180.60	21.26
1	133.32	260.14	185.60
2	559.97	2.84	111.76
3	391.43	351.05	76.39
4	394.78	352.61	194.35

	Water Mix (ml)	Plasticizer (gm)	Moderate Aggregator	Refined
0	201.66	16.11	1151.17	
1	175.99	6.27	1090.57	
2	295.23	11.95	1024.93	
3	299.14	19.00	1134.88	
4	235.54	17.02	1098.24	

	Formulation Duration (hrs)	Compression Strength MPa
0	344.43	79.89
1	28.86	59.80
2	237.68	77.86
3	208.81	71.74
4	266.84	76.07

```
dataset.isnull().sum()
```

Material Quantity (gm)	109
Additive Catalyst (gm)	109
Ash Component (gm)	109
Water Mix (ml)	109
Plasticizer (gm)	109
Moderate Aggregator	109
Refined Aggregator	109
Formulation Duration (hrs)	109

```
Compression Strength MPa      0
dtype: int64
```

NOTE: The below code eliminates all the null values in integer, strings format and by default fills the missing values with mean.

code reference - stack overflow

```
def clean_dataset(dataset):
    assert isinstance(dataset, pd.DataFrame)
    dataset.dropna(inplace=True)
    indices_to_keep = ~dataset.isin([np.nan, np.inf, -
np.inf]).any(axis=1)
    return dataset[indices_to_keep].astype(np.float64)
```

```
clean_dataset(dataset)
```

	Material Quantity (gm)	Additive Catalyst (gm)	Ash Component
0	486.42	180.60	
21.26			
1	133.32	260.14	
185.60			
2	559.97	2.84	
111.76			
3	391.43	351.05	
76.39			
4	394.78	352.61	
194.35			
...
6134	188.78	162.30	
142.65			
6135	349.87	291.45	
77.82			
6136	358.29	22.70	
17.99			
6137	445.25	275.59	
178.86			
6138	560.23	266.56	
167.14			

	Water Mix (ml)	Plasticizer (gm)	Moderate Aggregator	\
0	201.66	16.11	1151.17	
1	175.99	6.27	1090.57	
2	295.23	11.95	1024.93	
3	299.14	19.00	1134.88	
4	235.54	17.02	1098.24	
...	
6134	163.66	15.98	1003.82	

6135	188.26	25.82	925.10
6136	208.58	34.91	1081.07
6137	191.77	18.07	865.15
6138	175.49	10.63	1165.87

	Refined Aggregator	Formulation Duration (hrs)	Compression Strength MPa
0	708.50	344.43	
79.89			
1	1010.25	28.86	
59.80			
2	810.69	237.68	
77.86			
3	881.34	208.81	
71.74			
4	781.01	266.84	
76.07			
...	
...			
6134	1002.47	357.91	
50.61			
6135	1005.31	104.20	
54.24			
6136	792.44	302.76	
56.57			
6137	833.10	374.63	
58.21			
6138	894.53	360.96	
58.96			

[6030 rows x 9 columns]

dataset.isnull().sum()

Material Quantity (gm)	0
Additive Catalyst (gm)	0
Ash Component (gm)	0
Water Mix (ml)	0
Plasticizer (gm)	0
Moderate Aggregator	0
Refined Aggregator	0
Formulation Duration (hrs)	0
Compression Strength MPa	0

dtype: int64

OBSERVATION:

A clean dataset is achieved.

```
#scaling to increase the efficiency of the model
```

```
from sklearn.preprocessing import RobustScaler
```

```
transformer = RobustScaler().fit_transform(dataset)
transformer
```

```
array([[ 0.42668571, -0.05108089, -0.72481203, ..., -0.54009872,
         0.89630627,  0.88356314],
       [-0.95535094,  0.29558926,  0.54251012, ...,  1.06146171,
        -0.66358547,  0.00415846],
       [ 0.71456109, -0.82583682, -0.02691344, ...,  0.00228226,
         0.36863113,  0.79470344],
       ...,
       [-0.07481629, -0.73927824, -0.75002892, ..., -0.09458097,
         0.6903276 , -0.13722915],
       [ 0.26554596,  0.36292713,  0.49053403, ...,  0.12122499,
         1.04558767, -0.06544102],
       [ 0.71557874,  0.32357043,  0.40015423, ...,  0.44726925,
         0.9780156 , -0.03261107]])
```

```
scaled_dataset = pd.DataFrame(transformer, columns = dataset.columns)
scaled_dataset
```

	Material Quantity (gm)	Additive Catalyst (gm)	Ash Component
0	0.426686	-0.051081	-
0.724812			
1	-0.955351	0.295589	
0.542510			
2	0.714561	-0.825837	-
0.026913			
3	0.054894	0.691815	-
0.299672			
4	0.068006	0.698614	
0.609987			
...
...			
6025	-0.738280	-0.130840	
0.211297			
6026	-0.107772	0.432052	-
0.288645			
6027	-0.074816	-0.739278	-
0.750029			
6028	0.265546	0.362927	
0.490534			
6029	0.715579	0.323570	
0.400154			
	Water Mix (ml)	Plasticizer (gm)	Moderate Aggregator \
0	-0.358485	-0.011902	0.949160

1	-0.741276	-0.510256	0.573672
2	1.036833	-0.222588	0.166956
3	1.095139	0.134464	0.848225
4	0.146734	0.034186	0.621197
...
6025	-0.925142	-0.018486	0.036155
6026	-0.558306	0.479868	-0.451608
6027	-0.255294	0.940238	0.514809
6028	-0.505965	0.087364	-0.823068
6029	-0.748732	-0.289440	1.040244

	Refined Aggregator	Formulation Duration (hrs)	Compression Strength MPa
--	--------------------	----------------------------	--------------------------

0	-0.540099	0.896306	
---	-----------	----------	--

0.883563

1	1.061462	-0.663585	
---	----------	-----------	--

0.004158

2	0.002282	0.368631	
---	----------	----------	--

0.794703

3	0.377262	0.225924	
---	----------	----------	--

0.526811

4	-0.155247	0.512772	
---	-----------	----------	--

0.716349

...	
-----	-----	-----	--

...			
-----	--	--	--

6025	1.020169	0.962939	-
------	----------	----------	---

0.398118

6026	1.035242	-0.291173	-
------	----------	-----------	---

0.239221

6027	-0.094581	0.690328	-
------	-----------	----------	---

0.137229

6028	0.121225	1.045588	-
------	----------	----------	---

0.065441

6029	0.447269	0.978016	-
------	----------	----------	---

0.032611

[6030 rows x 9 columns]

```
from sklearn.model_selection import train_test_split
```

```
x = scaled_dataset[['Material Quantity (gm)', 'Additive Catalyst (gm)', 'Ash Component (gm)', 'Water Mix (ml)', 'Plasticizer (gm)', 'Moderate Aggregator', 'Refined Aggregator', 'Formulation Duration (hrs)']].values
```

```
y = scaled_dataset['Compression Strength MPa'].values
```

```
#hyper-parameter used here is random_state
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 150)
```

OBSERVATION:

After tuning the combinations of test_size and random_state continuously, highest r2_score was achieved at test_size = 0.2, random_state = 150

```
from sklearn.ensemble import RandomForestRegressor
#hyper-parameters n_estimators is tuned manually
regf = RandomForestRegressor(n_estimators=178)
regf.fit(x_train, y_train)
y_pred = regf.predict(x_test)

from sklearn.metrics import r2_score

r2_score(y_test, y_pred)
0.4632811125970906

r2_score(y_test, y_pred)*100
46.32811125970906
```